

# \*\*\* Applied Machine Learning Fundamentals \*\*\*

## *k*-Nearest Neighbors

Daniel Wehner

SAP SE

October 31, 2019



Find all slides on [GitHub](#)

# Lecture Overview

|                |                                |
|----------------|--------------------------------|
| Unit I         | Machine Learning Introduction  |
| Unit II        | Mathematical Foundations       |
| Unit III       | Bayesian Decision Theory       |
| Unit IV        | Probability Density Estimation |
| Unit V         | Regression                     |
| <b>Unit VI</b> | <b>Classification I</b>        |
| Unit VII       | Evaluation                     |
| Unit VIII      | Classification II              |
| Unit IX        | Clustering                     |
| Unit X         | Dimensionality Reduction       |

# Agenda October 31, 2019

## ① Introduction

- Overview of the Algorithm
- Derivation of the Algorithm

## ② Distance Metrics

- Properties of Distance Metrics
- Minkowski, Manhattan, Euclidean
- Cosine Similarity

## ③ $k$ -nearest Neighbors Algorithm

- General Procedure
- Calculation of Distances

Prediction of the Class Label

## ④ Choice of $k$

- Danger of Overfitting
- Selection Strategies

## ⑤ Wrap-Up

- Summary
- Self-Test Questions
- Lecture Outlook
- Recommended Literature and further Reading
- Meme of the Day


Section:  
**Introduction**



# Introduction

- **Basic idea:** Predict the class label based on nearby known examples
- Instance-based learning, a. k. a. **lazy learning**

We do not learn any model,  
the data speaks for itself!





# Derivation of the Algorithm

- Unconditional density:

$$p(\mathbf{x}) = \frac{k}{n \cdot v}$$

Remember non-parametric density estimation?

- Class priors:

$$p(\mathcal{C}_j) = \frac{n_j}{n}$$

Combine using Bayes' theorem:

$$p(\mathcal{C}_j|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_j) \cdot p(\mathcal{C}_j)}{p(\mathbf{x})} = \frac{\frac{k_j}{n_j \cdot v} \cdot \frac{n_j}{n}}{\frac{k}{n \cdot v}} = \frac{k_j}{k} \quad (1)$$

Section:  
**Distance Metrics**



# Distance Metrics

- How to measure the distance between two data points  $i$  and  $j$ ?  
 $\Rightarrow$  **distance metrics**
- Let  $d$  be a function  $d : (u, v) \mapsto \mathbb{R}^+$  (including 0)
- Function  $d$  has the following properties:

- ①  $d(u, v) = d(v, u)$  (**commutativity**)
- ②  $d(u, v) = 0 \Rightarrow u = v$
- ③  $d(u, k) \leq d(u, v) + d(v, k)$  (**triangle inequality**)



## Distance Metrics (Ctd.)

Minkowski distance:

$$d_p(u, v) = \left( \sum_{j=1}^m |x_j^{(u)} - x_j^{(v)}| \right)^{1/p} \quad (2)$$


Manhattan distance: ( $p = 1$ )

$$d_1(u, v) = \sum_{j=1}^m |x_j^{(u)} - x_j^{(v)}|$$

Euclidean distance: ( $p = 2$ )

$$d_2(u, v) = \sqrt{\sum_{j=1}^m |x_j^{(u)} - x_j^{(v)}|^2}$$

## Distance Metrics (Ctd.)



# Cosine Similarity


- **Similarity metrics** are an alternative to distance metrics
- **Example: Cosine similarity**
- The cosine similarity of two vectors  $\mathbf{a}$  and  $\mathbf{b}$  is the cosine of the angle:

$$\cos \angle(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\|_2 \cdot \|\mathbf{b}\|_2} = \frac{\sum_{j=1}^m a_j \cdot b_j}{\sqrt{\sum_{j=1}^m (a_j)^2} \cdot \sqrt{\sum_{j=1}^m (b_j)^2}} \quad (3)$$

- The dot product is defined as:

$$\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\|_2 \cdot \|\mathbf{b}\|_2 \cdot \cos \angle(\mathbf{a}, \mathbf{b}) \quad (4)$$

## Cosine Similarity (Ctd.)



Section:  
*k*-nearest Neighbors Algorithm



## Prediction with $k$ -Nearest Neighbors (Ctd.)

### $k$ -nearest neighbors algorithm:

- 1 Calculate the distances between the new data point and **all data points in the data set**
- 2 Sort the data points by distances **in ascending order**  
*(if similarity metrics are used, sort in descending order)*
- 3 Look at the first  $k$  examples and **count how often each class occurs**
- 4 Predict the class with **the maximum score**

# ① Calculation of Distances

| $v$      | $x_1$    | $x_2$    | $\mathcal{C}$ | $d_2(u, v)$ |
|----------|----------|----------|---------------|-------------|
| 1        | 0.66     | 0.24     | 1             | 0.23        |
| 2        | 0.25     | 0.79     | 1             | 0.67        |
| 3        | 0.16     | 0.81     | 1             | 0.73        |
| 4        | 0.57     | 0.21     | 1             | 0.13        |
| 5        | 0.21     | 0.72     | 1             | 0.62        |
| 6        | 0.66     | 0.27     | 1             | 0.24        |
| 7        | 0.27     | 0.11     | 0             | 0.19        |
| 8        | 0.39     | 0.13     | 0             | 0.07        |
| 9        | 0.39     | 0.86     | 0             | 0.71        |
| 10       | 0.44     | 0.67     | 0             | 0.52        |
| 11       | 0.31     | 0.33     | 0             | 0.23        |
| 12       | 0.03     | 0.51     | 0             | 0.55        |
| $\vdots$ | $\vdots$ | $\vdots$ | $\vdots$      | $\vdots$    |

- $\mathbf{x}^{(u)} = (0.45, 0.15)$
- Calculate the **Euclidean distance** between  $\mathbf{x}^{(u)}$  and all other data points  $\mathbf{x}^{(v)}$

Prediction is expensive!

## ②/③/④ Prediction of the Class Label

- Let  $k$  be set to 10
- Step ②: Sort data set by distances  
(cf. table)
- Step ③: Count the classes
  - Class 0: 3
  - Class 1: 7
- Step ④: Predict class 1!

| $x_1$    | $x_2$    | $\mathcal{C}$ | $d_2(u, v)$ |
|----------|----------|---------------|-------------|
| 0.51     | 0.17     | 1             | 0.06        |
| 0.39     | 0.13     | 0             | 0.07        |
| 0.52     | 0.17     | 1             | 0.08        |
| 0.43     | 0.23     | 0             | 0.08        |
| 0.47     | 0.03     | 0             | 0.12        |
| 0.52     | 0.26     | 1             | 0.13        |
| 0.57     | 0.21     | 1             | 0.13        |
| 0.53     | 0.25     | 1             | 0.13        |
| 0.58     | 0.12     | 1             | 0.14        |
| 0.59     | 0.13     | 1             | 0.14        |
| $\vdots$ | $\vdots$ | $\vdots$      | $\vdots$    |




Section:  
Choice of  $k$




# How to choose $k$ ?

The choice of  $k$  is important:



$k = 1$  (💀 overfitting 💀)



$k = 30$  (about right)

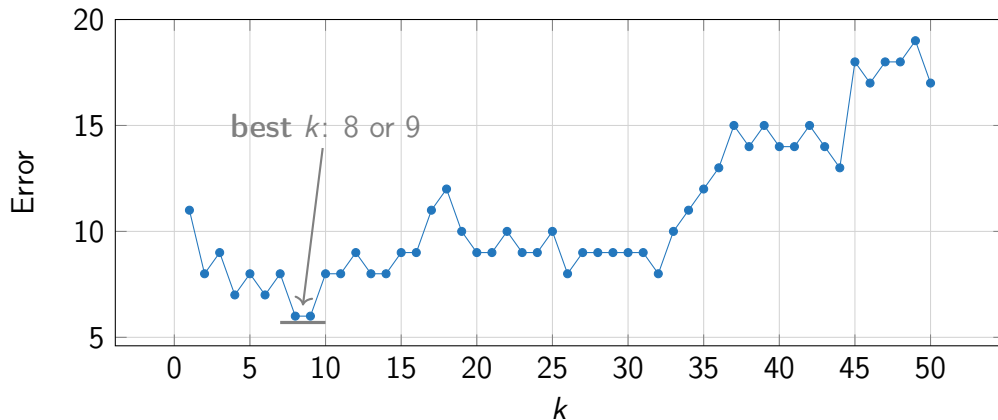
## How to choose *k*? (Ctd.)

- First of all, it is recommended to use **odd values** for *k*  
(*no tie-breaking necessary*)
- Compute *k* depending on the size of the data set  $\mathcal{D}$ :

$$k = \sqrt{\frac{n}{2}} \quad \text{or} \quad k = \sqrt{n} \quad (5)$$

- **Other strategy:** Evaluate different *k* on a dev set and choose the best one

## How to choose $k$ ? (Ctd.)



Section:  
**Wrap-Up**



# Summary

- The basic idea is to classify unknown instances **based on nearby examples**
- The algorithm is an example for **instance-based learning**
- **Distance metrics** allow to calculate the distance between data points:
  - Manhattan distance
  - Euclidean distance
  - Cosine similarity
- Choose the value of  $k$  wisely:
  - Too small: **Overfitting**
  - Too large: **Underfitting**



# Self-Test Questions

- 1 Outline the  $k$ -nearest neighbors algorithm.
- 2 What is instance-based learning (in contrast to model-based learning)?
- 3 How can you compute distances? What properties do distance metrics have?
- 4 What is the intuition behind the triangle inequality?
- 5 How can you choose  $k$ ?
- 6 Suppose you have a data set comprising  $n = 50$  examples.  
If you set  $k = n$ , what class does the algorithm predict?
- 7 What are advantages and disadvantages of the algorithm?

# What's next...?

|                |                                |
|----------------|--------------------------------|
| Unit I         | Machine Learning Introduction  |
| Unit II        | Mathematical Foundations       |
| Unit III       | Bayesian Decision Theory       |
| Unit IV        | Probability Density Estimation |
| Unit V         | Regression                     |
| <b>Unit VI</b> | <b>Classification I</b>        |
| Unit VII       | Evaluation                     |
| Unit VIII      | Classification II              |
| Unit IX        | Clustering                     |
| Unit X         | Dimensionality Reduction       |



# Recommended Literature and further Reading

## Meme of the Day



Thank you very much for the attention!

**Topic:** \*\*\* Applied Machine Learning Fundamentals \*\*\* *k*-Nearest Neighbors

**Date:** October 31, 2019

**Contact:**

Daniel Wehner

SAP SE

[daniel.wehner@sap.com](mailto:daniel.wehner@sap.com)

Do you have any questions?