*** Applied Machine Learning Fundamentals ***
Logistic Regression

Daniel Wehner, M.Sc.

SAP SE / DHBW Mannheim

Winter term 2020/2021

Lecture Overview

Unit I	Machine Learning Introduction
Unit II	Mathematical Foundations
Unit III	Bayesian Decision Theory
Unit IV	Probability Density Estimation
Unit V	Regression
Unit VI	**Classification I**
Unit VII	Evaluation
Unit VIII	Classification II
Unit IX	Clustering
Unit X	Dimensionality Reduction

Agenda for this Unit

Section:

Introduction

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

What is logistic Regression?
Why you should not use linear Regression

What is logistic Regression?

- Learning algorithm for **classification** *(despite the name...)*
- In its standard form it's applicable to **binary classification problems only**, but you can use techniques like:
  - **One-vs-One (OVO)**
  - **One-vs-Rest (OVR)**
- **Class labels:**
  - The 'positive class' is encoded as **1** / $\oplus$
  - The 'negative class' as **0** / $\ominus$
- **Probabilistic interpretation:** The output of the algorithm is between 0 and 1 *(probability of the instance belonging to the positive class)*

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

What is logistic Regression?
Why you should not use linear Regression

Why you should not use linear Regression...

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

What is logistic Regression?
Why you should not use linear Regression

Why you should not use linear Regression...

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

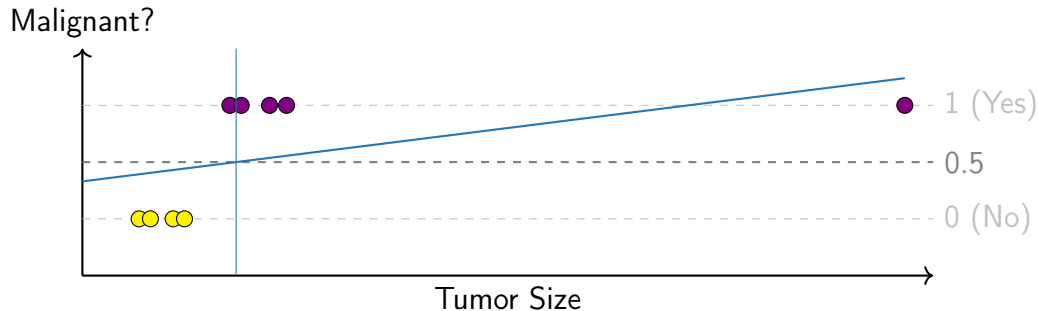What is logistic Regression?
Why you should not use linear Regression

Why you should not use linear Regression... (Ctd.)

- Linear regression: $h_{\theta}(x) = \theta^{\mathsf{T}} x$
- By putting a **threshold** at 0.5, we can turn linear regression into a classifier
  - If $h_{\theta}(x) \geqslant 0.5$, predict $y = 1$
  - If $h_{\theta}(x) < 0.5$, predict $y = 0$
- **Outliers affect the decision boundary**
- Furthermore, we only want $0 \leqslant h_{\theta}(x) \leqslant 1$
- Linear regression can output $h_{\theta}(x) \ll 0$ or $h_{\theta}(x) \gg 1$
- **We need a better strategy!**

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Important

Logistic Regression Model

- Remember that we want: $0 \leqslant h_{\boldsymbol{\theta}}(\boldsymbol{x}) \leqslant 1$

- **Solution:** Logistic / Sigmoid function:

$$g(z) = \frac{1}{1 + e^{-z}} \tag{1}$$

- We plug $\boldsymbol{\theta}^{\intercal}\boldsymbol{x}$ into the sigmoid function:

$$h_{\boldsymbol{\theta}}(\boldsymbol{x}) = g(\boldsymbol{\theta}^{\intercal}\boldsymbol{x}) = \frac{1}{1 + e^{-(\boldsymbol{\theta}^{\intercal}\boldsymbol{x})}} \tag{2}$$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Logistic/Sigmoid Function
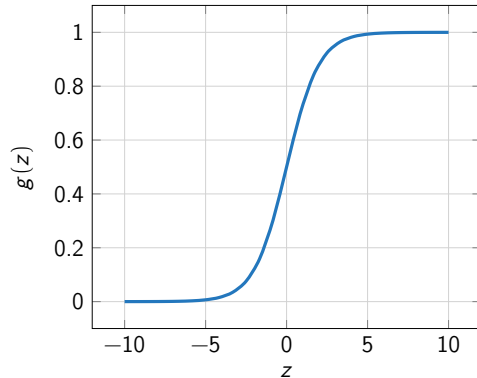
- $g(z)$ is symmetric around $z = 0$
- $0 \leqslant g(z) \leqslant 1$ holds true

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Where does the Sigmoid come from?

$$p(\mathcal{C}_1|\boldsymbol{x}) = \frac{p(\boldsymbol{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\boldsymbol{x})} = \frac{p(\boldsymbol{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{\sum_j p(\boldsymbol{x}, \mathcal{C}_j)} = \frac{p(\boldsymbol{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{\sum_j p(\boldsymbol{x}|\mathcal{C}_j)p(\mathcal{C}_j)} \tag{3}$$

$$= \frac{p(\boldsymbol{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\boldsymbol{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\boldsymbol{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \tag{4}$$

$$= \frac{1}{1 + p(\boldsymbol{x}|\mathcal{C}_2)p(\mathcal{C}_2)/(p(\boldsymbol{x}|\mathcal{C}_1)p(\mathcal{C}_1))} \tag{5}$$

$$= \frac{1}{1 + \exp\{-z\}} = g(z) \qquad \longrightarrow \textbf{logistic sigmoid} \tag{6}$$

$$z = \log \frac{p(\boldsymbol{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\boldsymbol{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \qquad \longrightarrow \textbf{log odds} \tag{7}$$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Interpretation of Hypothesis Output

- $h_{\boldsymbol{\theta}}(\boldsymbol{x})$ is interpreted as the probability of instance $\boldsymbol{x}$ belonging to class $y = 1$
- **Example:**

$$\boldsymbol{x} = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ tumorSize \end{bmatrix} \tag{8}$$

- If $h_{\boldsymbol{\theta}}(\boldsymbol{x}) = 0.7$, we have to tell the patient that there is a **70 % chance** of the tumor being malignant $\Rightarrow p(y = 1 | \boldsymbol{x}, \boldsymbol{\theta})$
- **Binary case:** $p(y = 0 | \boldsymbol{x}, \boldsymbol{\theta}) = 1 - p(y = 1 | \boldsymbol{x}, \boldsymbol{\theta})$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Training Setup

- We have a labeled training set ($\Rightarrow$ **supervised learning**):

$$\mathcal{D} = \left\{ (\boldsymbol{x}^{(1)}, y^{(1)}), (\boldsymbol{x}^{(2)}, y^{(2)}), \dots, (\boldsymbol{x}^{(n)}, y^{(n)}) \right\} = \left\{ (\boldsymbol{x}^{(i)}, y^{(i)}) \right\}_{i=1}^{n} \quad (9)$$

- Each $\boldsymbol{x}$ is a vector of features:

$$\boldsymbol{x} = \begin{bmatrix} x_0 \\ \vdots \\ x_m \end{bmatrix} \in \mathbb{R}^{m+1} \quad \text{and} \quad x_0 = 1 \quad \text{and} \quad y \in \{0, 1\} \quad (10)$$

- **How to choose the parameters θ?**

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Logistic Regression Cost Function

- Gradient descent is performed in order to find the parameters $\boldsymbol{\theta}$
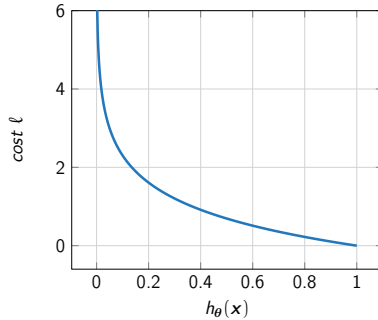- To this end, a cost function is needed:

$$\mathcal{J}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^{n} \ell(h_{\boldsymbol{\theta}}(\boldsymbol{x}^{(i)}), y^{(i)}) \qquad (11)$$

- The cost function $\ell(h_{\boldsymbol{\theta}}(\boldsymbol{x}), y)$ is defined as follows:
  *(square loss would be **non-convex...**)*

$$\ell(h_{\boldsymbol{\theta}}(\boldsymbol{x}), y) = \begin{cases} -\log(h_{\boldsymbol{\theta}}(\boldsymbol{x})) & \text{if } y = 1 \\ -\log(1 - h_{\boldsymbol{\theta}}(\boldsymbol{x})) & \text{if } y = 0 \end{cases} \qquad (12)$$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Logistic Regression Cost Function (Ctd.)

$y = 1$:

$y = 0$:

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary
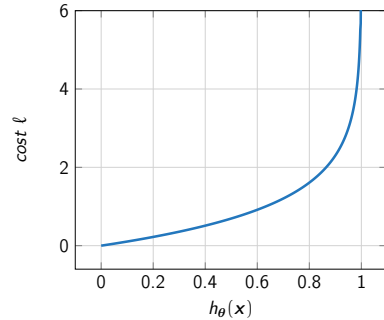
Logistic Regression Cost Function (Ctd.)

- $\ell(h_{\boldsymbol{\theta}}(\boldsymbol{x}), y)$ can be written in a more compact form:

$$\ell(h_{\boldsymbol{\theta}}(\boldsymbol{x}), y) = -y \log(h_{\boldsymbol{\theta}}(\boldsymbol{x})) - (1 - y) \log(1 - h_{\boldsymbol{\theta}}(\boldsymbol{x})) \qquad (13)$$

  - If $y = 1$, we get: $-\log(h_{\boldsymbol{\theta}}(\boldsymbol{x}))$
  - If $y = 0$, we get: $-\log(1 - h_{\boldsymbol{\theta}}(\boldsymbol{x}))$

- This gives the **cross entropy** cost function $\mathcal{J}(\boldsymbol{\theta})$:

$$\mathcal{J}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^{n} \left[-y^{(i)} \log(h_{\boldsymbol{\theta}}(\boldsymbol{x}^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\boldsymbol{\theta}}(\boldsymbol{x}^{(i)})) \right] \qquad (14)$$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Derivation of Cross Entropy

- The likelihood function can be written in the form:

$$\mathcal{L}(\boldsymbol{\theta}) = \prod_{i=1}^{n} h_{\boldsymbol{\theta}}(\boldsymbol{x}^{(i)})^{y^{(i)}} \cdot (1 - h_{\boldsymbol{\theta}}(\boldsymbol{x}^{(i)}))^{1-y^{(i)}} \tag{15}$$

- The cost function is then given by the **negative log-likelihood**:

$$\mathcal{J}(\boldsymbol{\theta}) = -\log \mathcal{L}(\boldsymbol{\theta}) \tag{16}$$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
**Model Training**
Decision Boundary

Gradient Descent

- The goal is to minimize $\mathcal{J}(\boldsymbol{\theta})$: $\boldsymbol{\theta}^* = \arg\min_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta})$

- Repeat until convergence {
    $$\boldsymbol{\theta}^{(t+1)} \longleftarrow \boldsymbol{\theta}^{(t)} - \alpha \nabla_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta}^{(t)}) \quad \textit{// simultaneously update all } \theta_j$$
  }

- The gradient $\nabla_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta})$ is given by:

$$\nabla_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^{n} \left(h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)} \right) \mathbf{x}^{(i)} \tag{17}$$

Algorithm looks identical to linear regression, but $h_{\theta}(x)$ is different!

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Decision Boundary

- **For classification we have to set a threshold**
- Suppose we predict $y = 1$, if $h_{\boldsymbol{\theta}}(\boldsymbol{x}) \geqslant 0.5$
  - This means $g(z) \geqslant 0.5$
  - This is equivalent to $z \geqslant 0$ and $\boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{x} \geqslant 0$
- Suppose we predict $y = 0$, if $h_{\boldsymbol{\theta}}(\boldsymbol{x}) < 0.5 \Rightarrow \boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{x} < 0$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Decision Boundary (Ctd.)

- Suppose we have the following hypothesis:
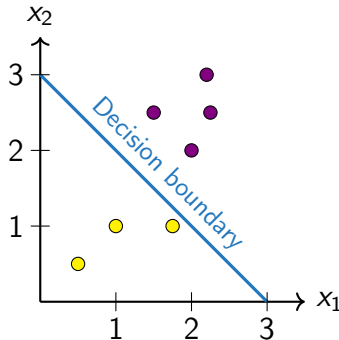
$$h_{\boldsymbol{\theta}}(\boldsymbol{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

- Using gradient descent we obtained the following coefficients:

$$\theta_0 = -3 \qquad \theta_1 = 1 \qquad \theta_2 = 1$$
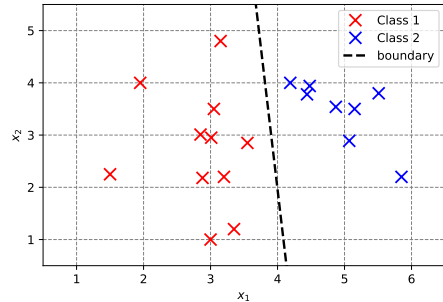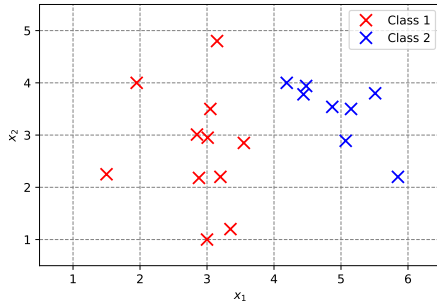
- Predict $y = 1$, if $-3 + x_1 + x_2 \geqslant 0$

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Decision Boundary (Ctd.)

- Predict $y = 1$, if $-3 + x_1 + x_2 \geqslant 0$
- The decision boundary satisfies $-3 + x_1 + x_2 = 0$
- If $x_2 = 0$, then $x_1 = 3$ and vice versa

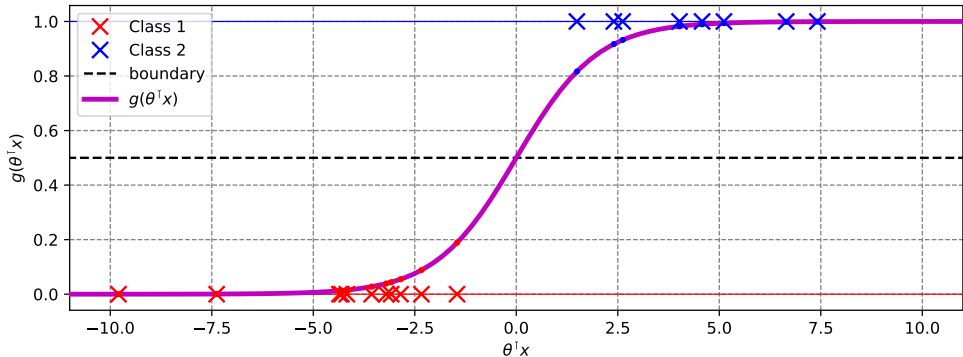**Logistic regression is not a maximum-margin classifier (although the cost function can be adjusted to get that $\Rightarrow$ Hinge loss)**

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Example: Decision Boundary

Where is the sigmoid function?

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Sigmoid Function
Probabilistic Interpretation
Model Training
Decision Boundary

Example: Logistic Function

Section:
Non-linear Data

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
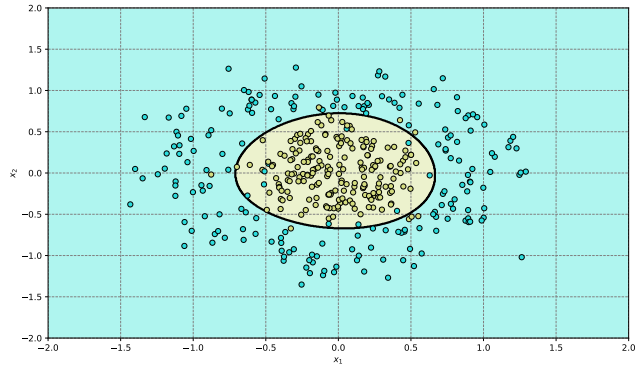Wrap-Up

Feature Mapping
Regularization

Non-Linear Decision Boundaries

- **Feature mapping** can be used to obtain non-linear decision boundaries
- **Example:**
  - Imagine a circular data set
  - Using the features...

  $$h_{\boldsymbol{\theta}}(\boldsymbol{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

  - ...the algorithm could e. g. choose: $\boldsymbol{\theta} = \begin{bmatrix} -1, 0, 0, 1, 1 \end{bmatrix}^{\mathsf{T}}$
  - So we would get: $x_1^2 + x_2^2 = 1 \Rightarrow$ **equation of a unit circle**

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Feature Mapping
Regularization

Example: Non-Linear Decision Boundary

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Feature Mapping
Regularization

Logistic Regression Cost Function (Ctd.)

- We should apply regularization for non-linear decision boundaries:

$$\frac{1}{n} \sum_{i=1}^{n} \left[-y^{(i)} \log(h_{\boldsymbol{\theta}}(\boldsymbol{x}^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\boldsymbol{\theta}}(\boldsymbol{x}^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^{m} \theta_j^2 \qquad (18)$$

- The last term prevents the parameters θ_j from becoming too large
- $\lambda \geqslant 0$ controls the degree of regularization
- This leads to smoother decision boundaries

Section:
Multi-Class Classification

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Multiple Classes
One-vs-Rest (OVR)
One-vs-One (OVO)

Multi-Class Classification

- Logistic regression can handle two classes only, namely **0** and **1**
- **What if there are more than two classes?**
- Two common techniques:
  - **One-vs-Rest (OVR)** $\Rightarrow$ One-against-All
  - **One-vs-One (OVO)** $\Rightarrow$ Pairwise classification
- Several classifiers are trained
- During prediction the classifiers **vote for the correct class**
- Such techniques can be used for all binary classifiers

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Multiple Classes
One-vs-Rest (OVR)
One-vs-One (OVO)

Multi-Class Classification: One-vs-Rest (OVR)

- **Train one classifier per class** (expert for that class)

- We get $|\mathcal{C}|$ classifiers

- The k-th classifier learns to distinguish the k-th class from all the others

- Set the labels of examples from class k to **1**, all the others to **0**

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Multiple Classes
One-vs-Rest (OVR)
One-vs-One (OVO)

Multi-Class Classification: One-vs-One (OVO)

- **Train one classifier for each pair of classes**
- We get $\binom{|\mathcal{C}|}{2}$ classifiers
- Ignore all other examples that do not belong to either of the two classes
- **Voting**: Count how often each class wins; the class with the highest score is predicted

Section:
Wrap-Up

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Summary
Self-Test Questions
Lecture Outlook
Recommended Literature and further Reading
Meme of the Day

Summary

- **Logistic regression is used for classification (!!!)**
- It is used for **binary classification problems** (generalizations exist)
- **Output:** Probability of instance belonging to positive class
- Apply a **threshold** to get the classification
- The algorithm minimizes the **cross entropy cost function**
- There is **no closed-form solution** (unlike for linear regression)
- **Basis functions** can be used for non-linear data
- **Multi-class classification:** One-vs-Rest, One-vs-One

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Summary
Self-Test Questions
Lecture Outlook
Recommended Literature and further Reading
Meme of the Day

Important

Self-Test Questions

1. Why should you not use linear regression for classification?
2. State the formula for the logistic function.
3. Why do we use cross entropy instead of the squared error?
4. Does logistic regression find the best-separating hyper-plane?
5. What techniques do you know for multi-class classification problems?

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Summary
Self-Test Questions
Lecture Outlook
Recommended Literature and further Reading
Meme of the Day

What's next...?

Unit I	Machine Learning Introduction
Unit II	Mathematical Foundations
Unit III	Bayesian Decision Theory
Unit IV	Probability Density Estimation
Unit V	Regression
Unit VI	**Classification I**
Unit VII	Evaluation
Unit VIII	Classification II
Unit IX	Clustering
Unit X	Dimensionality Reduction

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Summary
Self-Test Questions
Lecture Outlook
Recommended Literature and further Reading
Meme of the Day

Recommended Literature and further Reading I

📕 **[2] Pattern Recognition and Machine Learning**
Christopher Bishop. Springer. 2006.
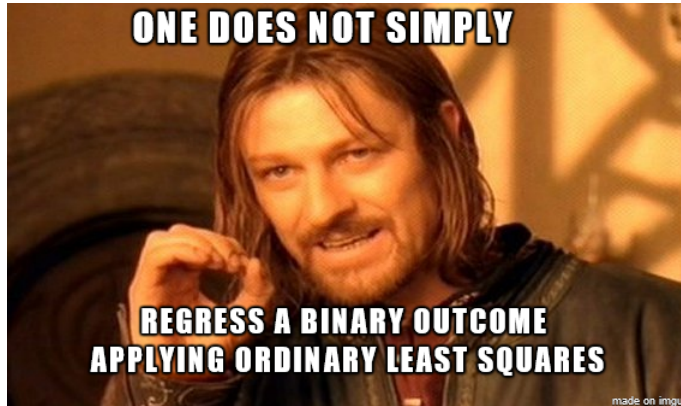→ <u>Link</u>, cf. chapter 4.3.2

📕 **[2] Machine Learning: A Probabilistic Perspective**
Kevin Murphy. MIT Press. 2012.
→ <u>Link</u>, cf. chapter 8

Introduction
Model Architecture
Non-linear Data
Multi-Class Classification
Wrap-Up

Summary
Self-Test Questions
Lecture Outlook
Recommended Literature and further Reading
Meme of the Day

Meme of the Day

Thank you very much for the attention!

Topic: *** Applied Machine Learning Fundamentals *** Logistic Regression
Term: Winter term 2020/2021

Contact:
Daniel Wehner, M.Sc.
SAP SE / DHBW Mannheim
daniel.wehner@sap.com

Do you have any questions?