

# \*\*\* Applied Machine Learning Fundamentals \*\*\*

## Logistic Regression

Daniel Wehner, M.Sc.

SAP SE / DHBW Mannheim

Winter term 2020/2021



Find all slides on [GitHub](#)

# Lecture Overview

Unit I	Machine Learning Introduction
Unit II	Mathematical Foundations
Unit III	Bayesian Decision Theory
Unit IV	Probability Density Estimation
Unit V	Regression
<b>Unit VI</b>	<b>Classification I</b>
Unit VII	Evaluation
Unit VIII	Classification II
Unit IX	Clustering
Unit X	Dimensionality Reduction

# Agenda for this Unit

## 1 Introduction

What is logistic Regression?  
Why you should not use linear Regression

## 2 Model Architecture

Sigmoid Function  
Probabilistic Interpretation  
Model Training  
Decision Boundary

## 3 Non-linear Data

Feature Mapping

Regularization

## 4 Multi-Class Classification

Multiple Classes  
One-vs-Rest (OVR)  
One-vs-One (OVO)

## 5 Wrap-Up

Summary  
Self-Test Questions  
Lecture Outlook  
Recommended Literature and further Reading  
Meme of the Day

Section:  
**Introduction**

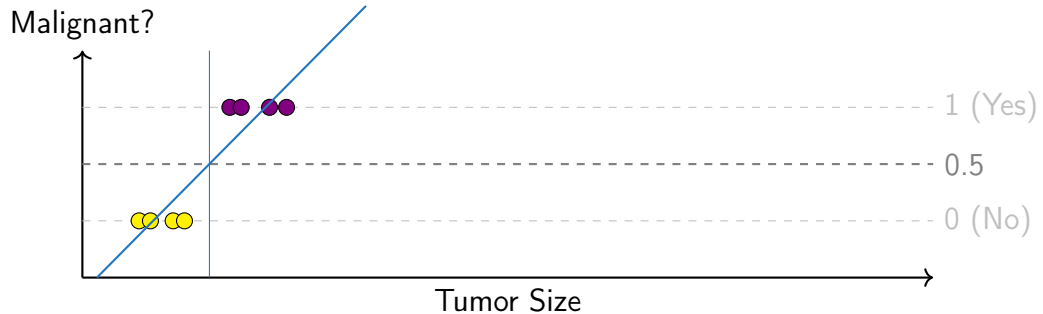


# What is logistic Regression?

- Learning algorithm for **classification** (*despite the name...*)
- In its standard form it's applicable to **binary classification problems only**, but you can use techniques like:
  - **One-vs-One (OVO)**
  - **One-vs-Rest (OVR)**
- **Class labels:**
  - The 'positive class' is encoded as  $1 / \oplus$
  - The 'negative class' as  $0 / \ominus$
- **Probabilistic interpretation:** The output of the algorithm is between 0 and 1 (*probability of the instance belonging to the positive class*)




# Why you should not use linear Regression...





# Why you should not use linear Regression...



## Why you should not use linear Regression... (Ctd.)

- Linear regression:  $h_{\theta}(\mathbf{x}) = \theta^T \mathbf{x}$
- By putting a **threshold** at 0.5, we can turn linear regression into a classifier
  - If  $h_{\theta}(\mathbf{x}) \geq 0.5$ , predict  $y = 1$
  - If  $h_{\theta}(\mathbf{x}) < 0.5$ , predict  $y = 0$
- **Outliers affect the decision boundary**
- Furthermore, we only want  $0 \leq h_{\theta}(\mathbf{x}) \leq 1$
- Linear regression can output  $h_{\theta}(\mathbf{x}) \ll 0$  or  $h_{\theta}(\mathbf{x}) \gg 1$
- **We need a better strategy!**



Section:  
**Model Architecture**





# Logistic Regression Model

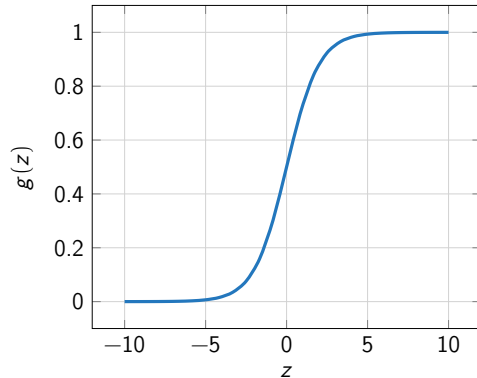
- Remember that we want:  $0 \leq h_{\theta}(\mathbf{x}) \leq 1$
- **Solution:** Logistic / Sigmoid function:

$$g(z) = \frac{1}{1 + e^{-z}} \quad (1)$$

- We plug  $\theta^T \mathbf{x}$  into the sigmoid function:

$$h_{\theta}(\mathbf{x}) = g(\theta^T \mathbf{x}) = \frac{1}{1 + e^{-(\theta^T \mathbf{x})}} \quad (2)$$

# Logistic/Sigmoid Function



- $g(z)$  is symmetric around  $z = 0$
- $0 \leq g(z) \leq 1$  holds true



## Where does the Sigmoid come from?

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x})} = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{\sum_j p(\mathbf{x}, \mathcal{C}_j)} = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{\sum_j p(\mathbf{x}|\mathcal{C}_j)p(\mathcal{C}_j)} \quad (3)$$

$$= \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (4)$$

$$= \frac{1}{1 + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)/(p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1))} \quad (5)$$

$$= \frac{1}{1 + \exp\{-z\}} = g(z) \quad \longrightarrow \text{logistic sigmoid} \quad (6)$$

$$z = \log \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad \longrightarrow \text{log odds} \quad (7)$$

# Interpretation of Hypothesis Output

- $h_{\theta}(\mathbf{x})$  is interpreted as the probability of instance  $\mathbf{x}$  belonging to class  $y = 1$
- **Example:**

$$\mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ tumorSize \end{bmatrix} \quad (8)$$

- If  $h_{\theta}(\mathbf{x}) = 0.7$ , we have to tell the patient that there is a **70 % chance** of the tumor being malignant  $\Rightarrow p(y = 1|\mathbf{x}, \theta)$
- **Binary case:**  $p(y = 0|\mathbf{x}, \theta) = 1 - p(y = 1|\mathbf{x}, \theta)$

# Training Setup

- We have a labeled training set ( $\Rightarrow$  **supervised learning**):

$$\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n \quad (9)$$

- Each  $\mathbf{x}$  is a vector of features:

$$\mathbf{x} = \begin{bmatrix} x_0 \\ \vdots \\ x_m \end{bmatrix} \in \mathbb{R}^{m+1} \quad \text{and} \quad x_0 = 1 \quad \text{and} \quad y \in \{0, 1\} \quad (10)$$

- How to choose the parameters  $\theta$ ?

# Logistic Regression Cost Function

- Gradient descent is performed in order to find the parameters  $\theta$
- To this end, a cost function is needed:

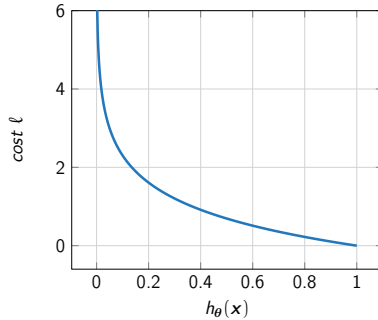
$$\mathcal{J}(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(h_{\theta}(\mathbf{x}^{(i)}), y^{(i)}) \quad (11)$$

- The cost function  $\ell(h_{\theta}(\mathbf{x}), y)$  is defined as follows:  
(square loss would be **non-convex...**)

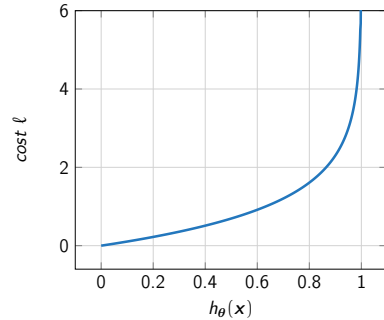
$$\ell(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases} \quad (12)$$

# Logistic Regression Cost Function (Ctd.)

$y = 1$ :



$y = 0$ :





# Logistic Regression Cost Function (Ctd.)

- $\ell(h_{\theta}(\mathbf{x}), y)$  can be written in a more compact form:

$$\ell(h_{\theta}(\mathbf{x}), y) = -y \log(h_{\theta}(\mathbf{x})) - (1 - y) \log(1 - h_{\theta}(\mathbf{x})) \quad (13)$$

- If  $y = 1$ , we get:  $-\log(h_{\theta}(\mathbf{x}))$
- If  $y = 0$ , we get:  $-\log(1 - h_{\theta}(\mathbf{x}))$
- This gives the **cross entropy** cost function  $\mathcal{J}(\theta)$ :

$$\mathcal{J}(\theta) = \frac{1}{n} \sum_{i=1}^n [-y^{(i)} \log(h_{\theta}(\mathbf{x}^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(\mathbf{x}^{(i)}))] \quad (14)$$



# Derivation of Cross Entropy

- The likelihood function can be written in the form:

$$\mathcal{L}(\boldsymbol{\theta}) = \prod_{i=1}^n h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)})^{y^{(i)}} \cdot (1 - h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}))^{1-y^{(i)}} \quad (15)$$

- The cost function is then given by the **negative log-likelihood**:

$$\mathcal{J}(\boldsymbol{\theta}) = -\log \mathcal{L}(\boldsymbol{\theta}) \quad (16)$$

# Gradient Descent


- The goal is to minimize  $\mathcal{J}(\boldsymbol{\theta})$ :  $\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta})$
- Repeat until convergence {  
 $\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} - \alpha \nabla_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta}^{(t)})$  // simultaneously update all  $\theta_j$   
}
- The gradient  $\nabla_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta})$  is given by:

$$\nabla_{\boldsymbol{\theta}} \mathcal{J}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \left( h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)} \right) \mathbf{x}^{(i)} \quad (17)$$

Algorithm looks identical to linear regression, but  $h_{\boldsymbol{\theta}}(\mathbf{x})$  is different!

# Decision Boundary

- For classification we have to set a threshold
- Suppose we predict  $y = 1$ , if  $h_{\theta}(\mathbf{x}) \geq 0.5$ 
  - This means  $g(z) \geq 0.5$
  - This is equivalent to  $z \geq 0$  and  $\theta^T \mathbf{x} \geq 0$
- Suppose we predict  $y = 0$ , if  $h_{\theta}(\mathbf{x}) < 0.5 \Rightarrow \theta^T \mathbf{x} < 0$



## Decision Boundary (Ctd.)

- Suppose we have the following hypothesis:

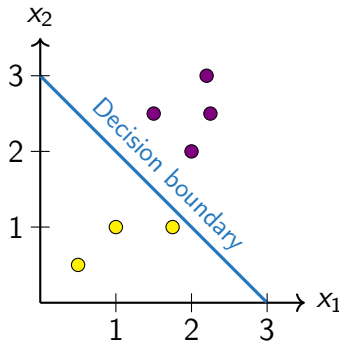
$$h_{\theta}(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

- Using gradient descent we obtained the following coefficients:

$$\theta_0 = -3 \quad \theta_1 = 1 \quad \theta_2 = 1$$

- Predict  $y = 1$ , if  $-3 + x_1 + x_2 \geq 0$

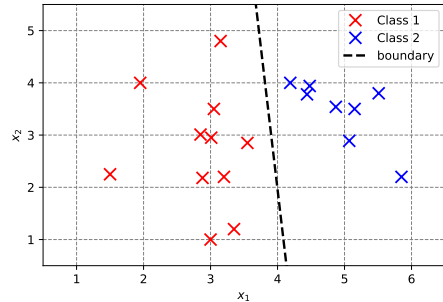
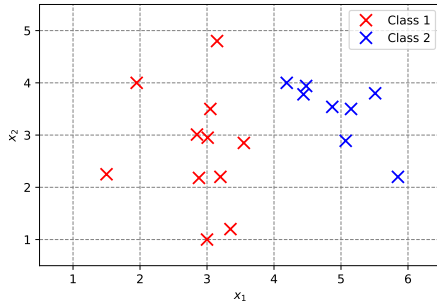
## Decision Boundary (Ctd.)



- Predict  $y = 1$ , if  $-3 + x_1 + x_2 \geq 0$
- The decision boundary satisfies  $-3 + x_1 + x_2 = 0$
- If  $x_2 = 0$ , then  $x_1 = 3$  and vice versa

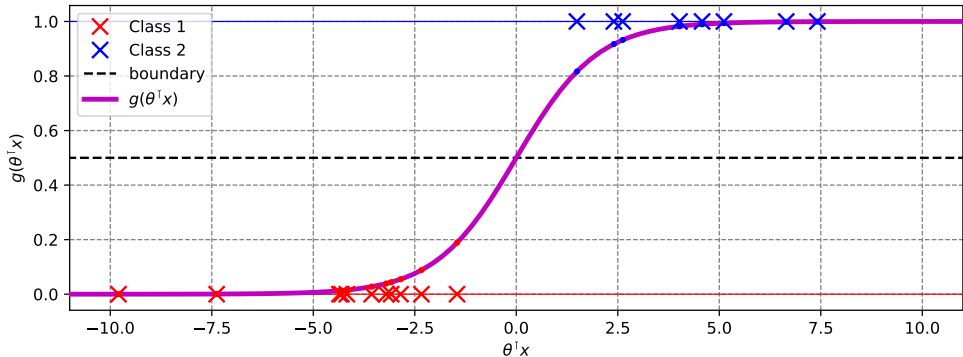
Logistic regression is not a maximum-margin classifier (although the cost function can be adjusted to get that  $\Rightarrow$  Hinge loss)

## Example: Decision Boundary



Where is the sigmoid function?

# Example: Logistic Function





Section:  
**Non-linear Data**



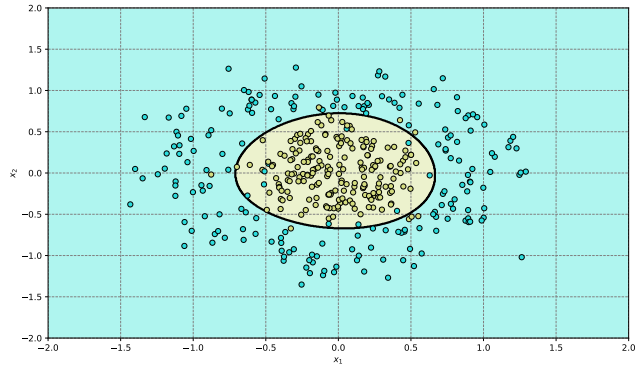
# Non-Linear Decision Boundaries

- **Feature mapping** can be used to obtain non-linear decision boundaries
- **Example:**
  - Imagine a circular data set
  - Using the features...

$$h_{\theta}(\mathbf{x}) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

- ...the algorithm could e. g. choose:  $\theta = [-1, 0, 0, 1, 1]^T$
- So we would get:  $x_1^2 + x_2^2 = 1 \Rightarrow$  **equation of a unit circle**

# Example: Non-Linear Decision Boundary



## Logistic Regression Cost Function (Ctd.)

- We should apply regularization for non-linear decision boundaries:

$$\frac{1}{n} \sum_{i=1}^n \left[ -y^{(i)} \log(h_{\theta}(\mathbf{x}^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(\mathbf{x}^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^m \theta_j^2 \quad (18)$$

- The last term prevents the parameters  $\theta_j$  from becoming too large
- $\lambda \geq 0$  controls the degree of regularization
- This leads to smoother decision boundaries

Section:  
**Multi-Class Classification**




# Multi-Class Classification


- Logistic regression can handle two classes only, namely 0 and 1
- **What if there are more than two classes?**
- Two common techniques:
  - **One-vs-Rest (OVR)**  $\Rightarrow$  One-against-All
  - **One-vs-One (OVO)**  $\Rightarrow$  Pairwise classification
- Several classifiers are trained
- During prediction the classifiers **vote for the correct class**
- Such techniques can be used for all binary classifiers

# Multi-Class Classification: One-vs-Rest (OVR)

- **Train one classifier per class**  
(expert for that class)
- We get  $|\mathcal{C}|$  classifiers
- The  $k$ -th classifier learns to distinguish the  $k$ -th class from all the others
- Set the labels of examples from class  $k$  to **1**, all the others to **0**



# Multi-Class Classification: One-vs-One (OVO)



- Train one classifier for each pair of classes
- We get  $\binom{|C|}{2}$  classifiers
- Ignore all other examples that do not belong to either of the two classes
- **Voting:** Count how often each class wins; the class with the highest score is predicted



Section:  
**Wrap-Up**



# Summary

- **Logistic regression is used for classification (!!!)**
- It is used for **binary classification problems** (generalizations exist)
- **Output:** Probability of instance belonging to positive class
- Apply a **threshold** to get the classification
- The algorithm minimizes the **cross entropy cost function**
- There is **no closed-form solution** (unlike for linear regression)
- **Basis functions** can be used for non-linear data
- **Multi-class classification:** One-vs-Rest, One-vs-One



# Self-Test Questions

- ① Why should you not use linear regression for classification?
- ② State the formula for the logistic function.
- ③ Why do we use cross entropy instead of the squared error?
- ④ Does logistic regression find the best-separating hyper-plane?
- ⑤ What techniques do you know for multi-class classification problems?

# What's next...?

Unit I	Machine Learning Introduction
Unit II	Mathematical Foundations
Unit III	Bayesian Decision Theory
Unit IV	Probability Density Estimation
Unit V	Regression
<b>Unit VI</b>	<b>Classification I</b>
Unit VII	Evaluation
Unit VIII	Classification II
Unit IX	Clustering
Unit X	Dimensionality Reduction

# Recommended Literature and further Reading I



## [1] Pattern Recognition and Machine Learning

*Christopher Bishop. Springer. 2006.*

→ [Link](#), cf. chapter 4.3.2

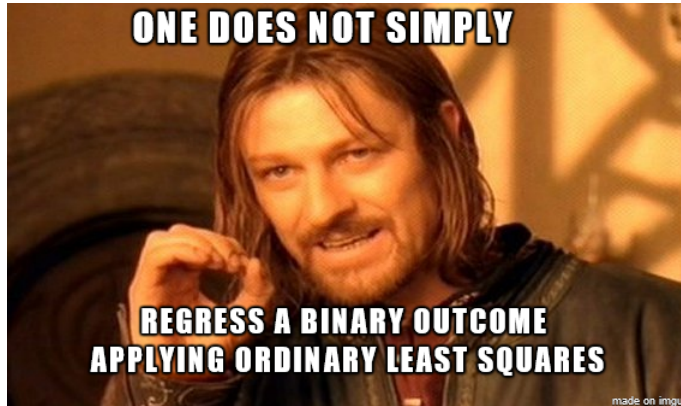


## [2] Machine Learning: A Probabilistic Perspective

*Kevin Murphy. MIT Press. 2012.*

→ [Link](#), cf. chapter 8

## Meme of the Day



Thank you very much for the attention!

**Topic:** \*\*\* Applied Machine Learning Fundamentals \*\*\* Logistic Regression

**Term:** Winter term 2020/2021

**Contact:**

Daniel Wehner, M.Sc.

SAP SE / DHBW Mannheim

[daniel.wehner@sap.com](mailto:daniel.wehner@sap.com)

Do you have any questions?