

*** Applied Machine Learning Fundamentals ***

Evaluation of ML Models

Daniel Wehner, M.Sc.

SAP SE / DHBW Mannheim

Winter term 2023/2024



Find all slides on [GitHub](#) (DaWe1992/Applied_ML_Fundamentals)

Lecture Overview

Unit I	Machine Learning Introduction
Unit II	Mathematical Foundations
Unit III	Bayesian Decision Theory
Unit IV	Regression
Unit V	Classification I
Unit VI	Evaluation
Unit VII	Classification II
Unit VIII	Clustering
Unit IX	Dimensionality Reduction

Agenda for this Unit

- 1 Evaluation Methods and Data Splits
- 2 Evaluation Metrics for Classifiers
- 3 Evaluation Metrics for Regressors
- 4 Model Selection and Model Complexity
- 5 Wrap-Up

Section: Evaluation Methods and Data Splits

Introduction
Out-of-Sample Testing and Data Splits
Cross-Validation / LOO-Validation

Evaluation of trained Models

- ① **Validation through experts:** A domain expert checks the plausibility
 - Subjective, time-intensive, and costly
 - Often the only option
- ② **Validation on data:** Evaluate the performance on a **separate (!)** test set
 - Labeled data is scarce and could be better used for training
 - Fast and simple, no domain knowledge needed
- ③ **On-line validation:** Test the model in a field test
 - Bad models may be costly (e. g. autonomous driving)
 - Gives the best estimate for the overall utility



Out-of-Sample Testing

- The performance cannot be measured on the training data (why?)
- Usually, a portion of the available data is reserved for testing
 - $\frac{2}{3}$ for training, $\frac{1}{3}$ for testing (evaluation)
 - The model is trained on the training set and evaluated on the test set
- **Problems:**
 - Waste of data
 - Labeling may be expensive
- **Solution:** **Cross-Validation (X-Val)**



Three Splits: Train, Dev/Validation, Test

In practice it is also common to split the data into three portions:

① **Training set** (used for training as before)

② **Dev/Validation set**

- Used for hyper-parameter tuning of the model
- Using the test set for that would be cheating

③ **Test set**

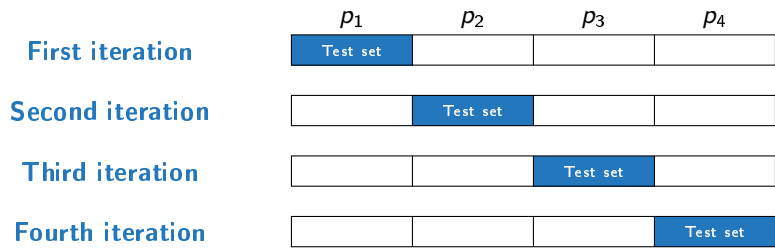
- The final model is tested on the test set
- The test set is used to estimate the **generalization error**

Stratified splits have the same class distribution as the entire dataset



Cross-Validation (k -fold X-Val)

- Split the dataset into k equally sized partitions $P = \{p_1, p_2, \dots, p_k\}$
- For each partition p_i : Use $P \setminus \{p_i\}$ for training and p_i for testing
- Average the results



Leave-One-Out Cross-Validation (LOO X-Val)

- n -fold X-Val
 - n is the number of examples
 - Use $n - 1$ examples for training, one example for testing
- **Properties:**
 - Makes best use of the data
 - Very expensive for large datasets (large n)

Note: We get k trained models when using k -fold X-Val!

- Which of these models is used in production?
- **Answer:** None of them. X-Val is only used for error estimation. The final model is trained on the entire dataset

Section: Evaluation Metrics for Classifiers

Confusion Matrices
Drawback of Accuracy
Precision, Recall and F1-Score
ROC and AUC

Types of Errors

- **Type I Error:** False negatives
 - An instance which is labeled \ominus is classified as \oplus
 - E. g. a spam e-mail is not detected
- **Type II Error:** False positives
 - An instance which is labeled \oplus is classified as \ominus
 - E. g. a non-spam (ham) e-mail is classified as spam

a. k. a. α/β error

Depending on the context the costs of false negatives and false positives can be different!



Confusion Matrices (two Classes)

- How often is class \mathcal{C}_i confused with class \mathcal{C}_j ?
- Calculate **accuracy**:

		Classified	
		\oplus	\ominus
Is	\oplus	#tp	#fn
	\ominus	#fp	#tn

$$accuracy = \frac{\#tp + \#tn}{\#tp + \#tn + \#fp + \#fn}$$

$$error = 1 - accuracy$$

tp true positive
fn false negative
fp false positive
tn true negative

Confusion Matrices (multiple Classes)

		Classified			
		A	B	C	Σ
Is	A	$n_{A,A}$	$n_{B,A}$	$n_{C,A}$	n_A
	B	$n_{A,B}$	$n_{B,B}$	$n_{C,B}$	n_B
	C	$n_{A,C}$	$n_{B,C}$	$n_{C,C}$	n_C
	Σ	$\overline{n_A}$	$\overline{n_B}$	$\overline{n_C}$	n

$$accuracy = \frac{n_{A,A} + n_{B,B} + n_{C,C}}{n}$$



Drawback of Accuracy

- Real-world datasets are usually **imbalanced**, i. e. some classes appear more frequently than others
- **Example:**
 - A dataset \mathcal{D} contains two classes \mathcal{C}_1 and \mathcal{C}_2
 - \mathcal{C}_1 appears 99 % of the time, \mathcal{C}_2 only 1 % of the time
 - It is easy to reach 99 % accuracy by always predicting the majority class
 - **Is this useful?** *Probably not...*

We need some more sophisticated evaluation metrics!



Precision and Recall

Precision: Ratio of #tp to all instances predicted as \oplus

$$P := \frac{\#tp}{\#tp + \#fp} \quad (1)$$

Recall (Sensitivity): Ratio of #tp to all instances actually labeled as \oplus

$$R := \frac{\#tp}{\#tp + \#fn} \quad (2)$$



Precision-Recall Trade-Off

There is a trade-off between precision and recall:

It is very easy to get 100 % precision:

- Simply classify one instance as \oplus where you are absolutely sure
- But recall is bad... (*many \oplus -instances are not detected*)

It is also quite easy to achieve 100 % recall:

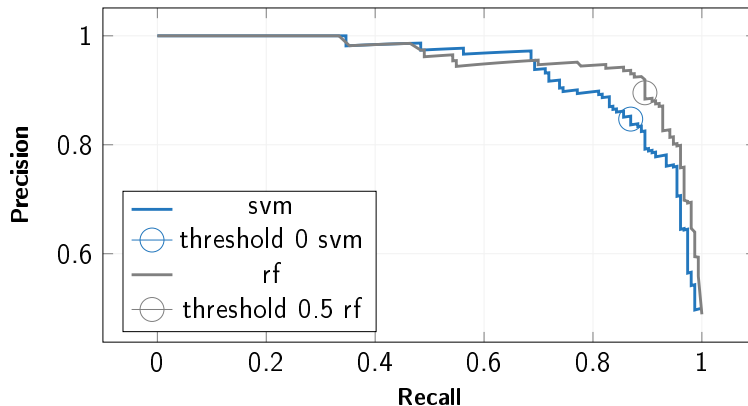
- Classify all instances as \oplus
- But precision is bad... (*many \ominus -instances are detected*)

Precision-Recall Curves / P-R-Curves

- Visualization of the precision-recall trade-off
- Influence precision and recall by changing thresholds
- **Example:**
 - Consider a ranker, e. g. a logistic regression classifier
 - It outputs probabilities for each class
 - The threshold when to predict \oplus can be changed
 - This has an influence on precision and recall

A P-R-curve plots precision and recall for all possible thresholds.

Precision-Recall Curves / P-R-Curves (Ctd.)



Combining Precision and Recall: F1-Score

- When to use precision, when recall?
- This depends on the cost of fp and fn
 - If fp are expensive \Rightarrow **use precision!**
 - If fn are expensive \Rightarrow **use recall!**
- **F1-score** (*harmonic mean of precision and recall*)

Why the harmonic mean?

$$F_1 := \frac{2 \cdot P \cdot R}{P + R} \quad F_\beta := (1 + \beta^2) \cdot \frac{P \cdot R}{(\beta^2 \cdot P) + R} \quad (\beta \in \mathbb{R}^+) \quad (3)$$

- Large β emphasizes recall



Calculation for multiple Classes (Example Precision)

- Precision must be calculated for each class separately
- For K classes we get K results. **How to combine the results?**
 - **Macro average:** Calculate P for each class and average the result

$$P_{macro} := \frac{1}{K} \sum_{k=1}^K P_k \quad (4)$$

- **Micro average:** Sum $\#tp$ and $\#fp$ for all classes and calculate P

$$P_{micro} := \sum_{k=1}^K \#tp_k / \sum_{k=1}^K (\#tp_k + \#fp_k) \quad (5)$$

Calculation for multiple Classes (Example Precision, Ctd.)

		Classified				
		A	B	C	D	Σ
Is	A	40	12	4	8	64
	B	7	51	2	0	60
	C	2	17	27	11	57
	D	39	4	15	8	66
	Σ	88	84	48	27	247

Calculation for multiple Classes (Example Precision, Ctd.)

$$P_A = \frac{40}{40 + 48} = 0.45$$

$$P_B = \frac{51}{51 + 33} = 0.61$$

$$P_C = \frac{27}{27 + 21} = 0.56$$

$$P_D = \frac{8}{19} = 0.30$$

$$P_{macro} = \frac{0.45 + 0.61 + 0.56 + 0.30}{4} = 0.48$$

$$P_{micro} = \frac{40 + 51 + 27 + 8}{(40 + 51 + 27 + 8) + (48 + 33 + 21 + 19)} = 0.51$$

ROC-Curves

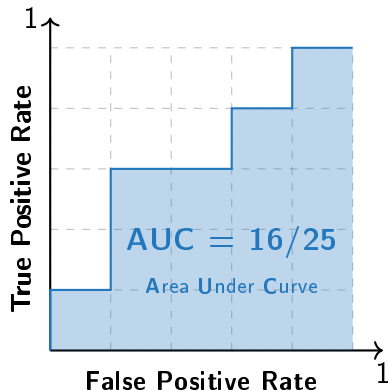
- ROC is short for **R**eciever **O**perating **C**haracteristic
- Borrowed from signal theory (*hence the name*)
- Uses *true positive rate* (recall) and *false positive rate* = $\frac{\#fp}{\#fp + \#tn}$

General procedure:

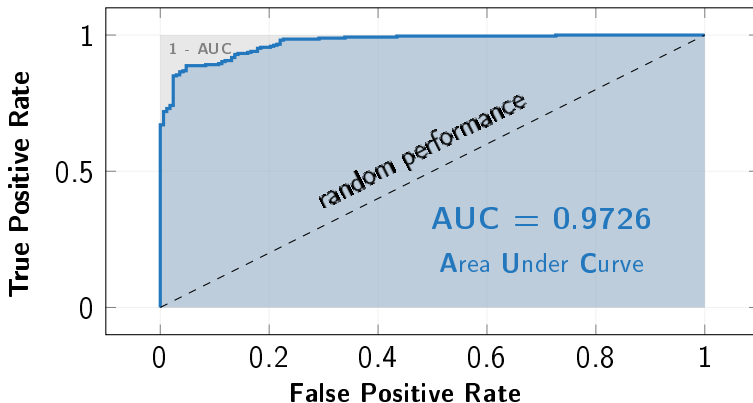
- Rank test instances by decreasing certainty of class \oplus
- Start at the origin (0,0)
- If the next instance in the ranking is \oplus : move $1/|\oplus|$ up
- If the next instance in the ranking is \ominus : move $1/|\ominus|$ right

Sample ROC-Curve (Example I)

Rank	Prob.	True class
1	0.95	\oplus
2	0.85	\ominus
3	0.78	\oplus
4	0.75	\oplus
5	0.62	\ominus
6	0.41	\ominus
7	0.37	\oplus
8	0.22	\ominus
9	0.15	\oplus
10	0.05	\ominus



Sample ROC-Curve (Example II)



ROC-Curve Interpretation

- AUC can be interpreted as the probability of a positive example always being listed before a negative example
- A high AUC value entails a good class separation:
 - AUC = 1.0:** All \oplus listed before all \ominus (desiderata)
 - AUC = 0.5:** Random ordering (worst case)
 - AUC = 0.0:** All \ominus listed before all \oplus (not the worst case \Rightarrow Invert classification)

Analogy: In a quiz you are allowed to answer those questions first where you feel the most confident (ranking). If you answer the first questions wrong, you won't perform very well overall \Rightarrow **small AUC**.

Section: Evaluation Metrics for Regressors

R^2 , RMSE and MAE
An Example

R^2 , RMSE and MAE

- Coefficient of determination R^2 :

$$R^2 := \frac{\sum_{i=1}^n (h_{\theta}(\mathbf{x}^{(i)}) - \bar{y})^2}{\sum_{i=1}^n (y^{(i)} - \bar{y})^2} = \frac{\text{Variance explained by model}}{\text{Total variance}} \quad R^2 \in [0, 1] \quad (6)$$

- Root mean square error (RMSE):

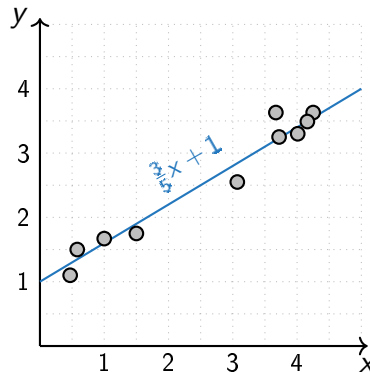
$$RMSE := \left(\frac{1}{n} \cdot \sum_{i=1}^n \left(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)} \right)^2 \right)^{1/2} \quad (7)$$

- Mean absolute error (MAE):

$$MAE := \frac{1}{n} \cdot \sum_{i=1}^n |h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}| \quad (8)$$

Evaluation of Regressors (Ctd.)

$x^{(i)}$	$y^{(i)}$	$h_{\theta}(x^{(i)})$
0.47	1.10	1.28
0.58	1.50	1.35
1.00	1.67	1.60
1.50	1.75	1.90
3.07	2.55	2.84
3.67	3.63	3.20
3.72	3.25	3.23
4.01	3.30	3.41
4.16	3.49	3.50
4.25	3.63	3.55
$\bar{y} = 2.59$		



Evaluation of Regressors (Ctd.)

- Coefficient of determination:

$$R^2 = \frac{(1.28 - 2.59)^2 + \dots + (3.55 - 2.59)^2}{(1.10 - 2.59)^2 + \dots + (3.63 - 2.59)^2} = \frac{7.97}{8.89} = 0.90 \quad (9)$$

- Root mean square error:

$$RMSE = \left(\frac{1}{10} \cdot [(1.28 - 1.10)^2 + \dots + (3.55 - 3.63)^2] \right)^{1/2} = 0.19 \quad (10)$$

- Mean absolute error:

$$MAE = \frac{1}{10} \cdot (|1.28 - 1.10| + \dots + |3.55 - 3.63|) = 0.15 \quad (11)$$

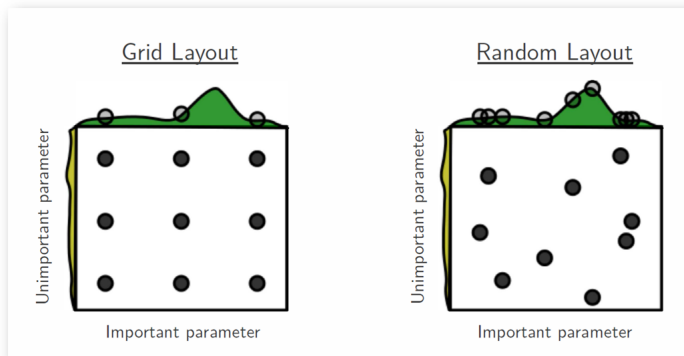
Section: Model Selection and Model Complexity

Hyper-Parameter Tuning: Grid Search and Random Search
Bias and Variance

Hyper-Parameter Tuning with Grid Search

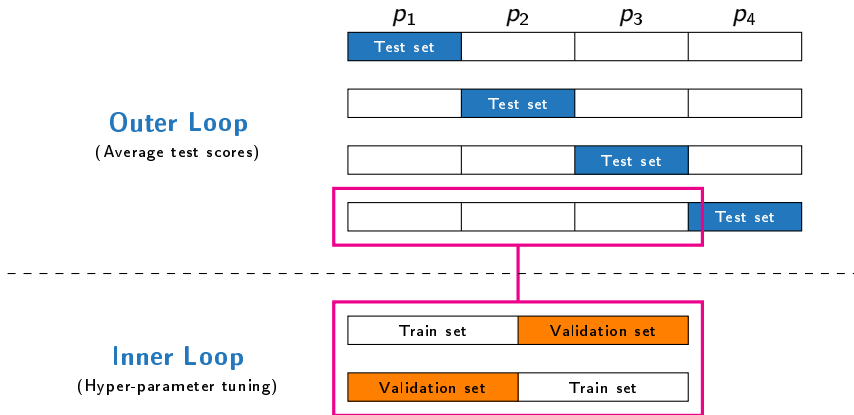
- **Grid search** is applied to find **optimal hyper-parameter settings**
- Hyper-parameter tuning should be done on the **dev** set
- We have to specify the search space / ranges of hyper-parameter values
- Grid search will try **all combinations** to find the best model
 - **Computationally very expensive**
 - Scikit-learn provides parameters to parallelize the search
(`n_jobs=-1` \Rightarrow use all cores available)
 - May not find the optimal setting \Rightarrow **Random search**

Grid Search vs. random Search



cf. Bergstra/Bengio.2012, <https://www.jmlr.org/papers/volume13/bergstra12a/bergstra12a.pdf>, page 284

Nested Cross-Validation



Bias-Variance Decomposition for MSE

- Suppose that we have a training set $\mathcal{D} = \{(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)})\}$
- There is an **unknown** function $f(x)$ such that

$$y^{(i)} = f(x^{(i)}) + \varepsilon \quad (12)$$

- The noise term ε has zero mean and variance σ_ε^2 : $\mathbb{E}\{\varepsilon\} = 0, \mathbb{V}\{\varepsilon\} = \sigma_\varepsilon^2$
- Based on the dataset \mathcal{D} we want to find a function $\hat{f}(x; \mathcal{D})$ which approximates $f(x)$ 'as well as possible'
- For this we try to minimize the mean squared error $\mathbb{E} \left\{ (y^{(i)} - \hat{f}(x^{(i)}))^2 \right\}$



Bias-Variance Decomposition for MSE (Ctd.)

- We can decompose the **expected error** of \hat{f} on an unseen sample x :

Bias-Variance decomposition of the mean squared error:

$$\mathbb{E}_{\mathcal{D}, \epsilon} \left\{ (y - \hat{f}(x; \mathcal{D}))^2 \right\} = \mathbb{B}_{\mathcal{D}}^2 \{ \hat{f}(x; \mathcal{D}) \} + \mathbb{V}_{\mathcal{D}} \{ \hat{f}(x; \mathcal{D}) \} + \sigma_{\epsilon}^2 \quad (13)$$

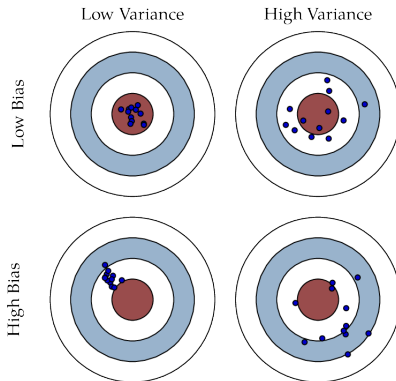
- $\mathbb{B}_{\mathcal{D}}$ is the **bias**, and $\mathbb{V}_{\mathcal{D}}$ the **variance** of the model \hat{f}
- σ_{ϵ}^2 is **irreducible**



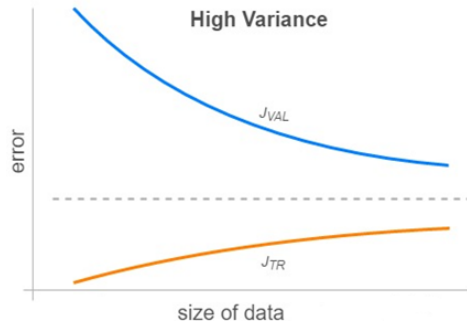
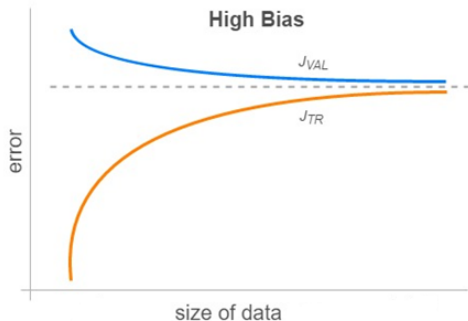
Bias and Variance

The **bias** results from erroneous assumptions in the learning algorithm. High bias can cause an algorithm to miss the relevant relations between features and target outputs (**underfitting**).

The **variance** is an error from sensitivity to small fluctuations in the training set. High variance may result from an algorithm modeling the random noise in the training data (**overfitting**).



Bias and Variance (Ctd.)





Recall: Expectation and Variance

- Let \mathcal{X} be a random variable. $\Omega(\mathcal{X})$ is the domain of \mathcal{X} (set of possible values \mathcal{X} can take)
- Definition of **expectation** (discrete case):

$$\mathbb{E}\{\mathcal{X}\} := \sum_{k \in \Omega(\mathcal{X})} k \cdot p(\mathcal{X} = k) \quad (14)$$

- Definition of **variance** (discrete case):

$$\mathbb{V}\{\mathcal{X}\} := \sum_{k \in \Omega(\mathcal{X})} (k - \mathbb{E}\{\mathcal{X}\})^2 \cdot p(\mathcal{X} = k) \quad (15)$$



Recall: Expectation and Variance (Ctd.)

Let \mathcal{X} and \mathcal{Y} be random variables and $a, b \in \mathbb{R}$:

- **Linearity** of \mathbb{E} (very important!):

$$\mathbb{E}\{a\mathcal{X} + b\mathcal{Y}\} = a\mathbb{E}\{\mathcal{X}\} + b\mathbb{E}\{\mathcal{Y}\} \quad (16)$$

- If \mathcal{X} and \mathcal{Y} are independent: $\mathbb{E}\{\mathcal{X}\mathcal{Y}\} = \mathbb{E}\{\mathcal{X}\}\mathbb{E}\{\mathcal{Y}\}$
- $\mathbb{V}\{\mathcal{X}\} = \mathbb{E}\{\mathcal{X} - \mathbb{E}\{\mathcal{X}\}\}^2 = \mathbb{E}\{\mathcal{X}^2\} - \mathbb{E}^2\{\mathcal{X}\}$
- \mathbb{V} is **not** linear: $\mathbb{V}\{a + b\mathcal{X}\} = b^2\mathbb{V}\{\mathcal{X}\}$
- However, if \mathcal{X} and \mathcal{Y} are uncorrelated: $\mathbb{V}\{\mathcal{X} + \mathcal{Y}\} = \mathbb{V}\{\mathcal{X}\} + \mathbb{V}\{\mathcal{Y}\}$



Derivation of the Bias-Variance Decomposition for MSE

- The MSE is given by

$$\begin{aligned}\text{MSE} &:= \mathbb{E}\{(y - \hat{f})^2\} = \mathbb{E}\{y^2 - 2y\hat{f} + \hat{f}^2\} \\ &= \underbrace{\mathbb{E}\{y^2\}}_{\text{①}} - 2 \underbrace{\mathbb{E}\{y\hat{f}\}}_{\text{②}} + \underbrace{\mathbb{E}\{\hat{f}^2\}}_{\text{③}}\end{aligned}\quad (17)$$

- Term ③ is straight-forward:

$$\begin{aligned}\mathbb{E}\{\hat{f}^2\} &= \mathbb{E}\{\hat{f}^2\} - \mathbb{E}^2\{\hat{f}\} + \mathbb{E}^2\{\hat{f}\} \\ &= \mathbb{V}\{\hat{f}\} + \mathbb{E}^2\{\hat{f}\}\end{aligned}\quad (18)$$



Derivation of the Bias-Variance Decomposition for MSE (Ctd.)

- We rewrite term ❶:

$$\begin{aligned}\mathbb{E}\{y^2\} &= \mathbb{E}\{(f + \varepsilon)^2\} && \text{Definition of } y \\ &= \mathbb{E}\{f^2\} + 2\mathbb{E}\{f\varepsilon\} + \mathbb{E}\{\varepsilon^2\} && \text{Linearity of } \mathbb{E} \\ &= f^2 + 2f\mathbb{E}\{\varepsilon\} + \mathbb{E}\{\varepsilon^2\} && f \text{ does not depend on } \mathcal{D} \\ &= f^2 + 2f \cdot 0 + \sigma_\varepsilon^2 && \mathbb{E}\{\varepsilon\} = 0 \quad (19)\end{aligned}$$



Derivation of the Bias-Variance Decomposition for MSE (Ctd.)

- We rewrite term ②:

$$\mathbb{E}\{y\hat{f}\} = \mathbb{E}\{(f + \varepsilon)\hat{f}\}$$

Definition of y

$$= f\mathbb{E}\{\hat{f}\} + \mathbb{E}\{\varepsilon\hat{f}\}$$

Linearity of \mathbb{E}

$$= f\mathbb{E}\{\hat{f}\} + \mathbb{E}\{\varepsilon\}\mathbb{E}\{\hat{f}\}$$

\hat{f} and ε are independent

$$= f\mathbb{E}\{\hat{f}\}$$

$$\mathbb{E}\{\varepsilon\} = 0 \quad (20)$$



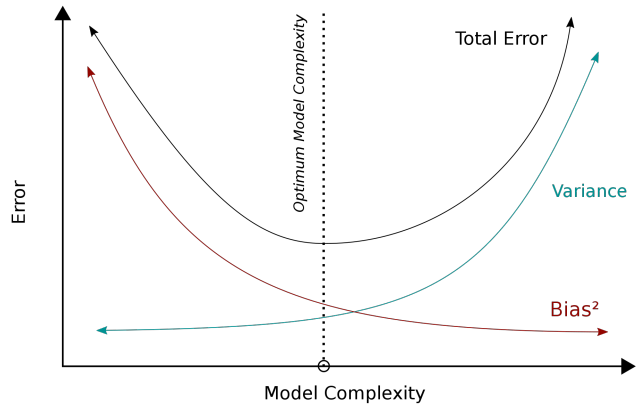
Derivation of the Bias-Variance Decomposition for MSE (Ctd.)

- Let us now plug these results into our MSE definition:

$$\begin{aligned}\text{MSE} &= \mathbb{E}\{y^2\} - 2\mathbb{E}\{y\hat{f}\} + \mathbb{E}\{\hat{f}^2\} = f^2 + \sigma_\varepsilon^2 - 2f\mathbb{E}\{\hat{f}\} + \mathbb{V}\{\hat{f}\} + \mathbb{E}^2\{\hat{f}\} \\ &= (f - \mathbb{E}\{\hat{f}\})^2 + \mathbb{V}\{\hat{f}\} + \sigma_\varepsilon^2 = \mathbb{E}^2\left\{(f - \hat{f})\right\} + \mathbb{V}\{\hat{f}\} + \sigma_\varepsilon^2 \\ &= \mathbb{B}^2\{\hat{f}\} + \mathbb{V}\{\hat{f}\} + \sigma_\varepsilon^2\end{aligned}\tag{21}$$



Bias, Variance, and Model Complexity

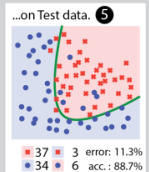


Use early stopping!

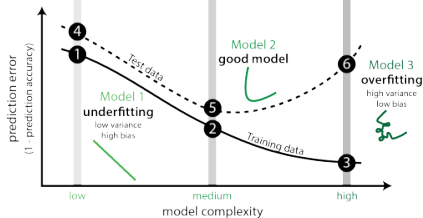
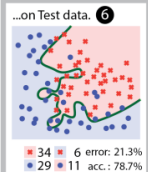
Model 1...



Model 2...



Model 3...



Section: Wrap-Up

Summary
Self-Test Questions
Lecture Outlook

Summary

- **Out-of-sample testing:** Split data into train, dev and test sets
- Cross-validation makes **maximum use of the data**
- Confusion matrices reveal **which classes are frequently confused**
- Precision, recall and F1 are **more robust w. r. t. imbalanced datasets**
- ROC curves are used for the evaluation of rankers
- Hyper-parameters are optimized using **grid search** or **random search**
- Keep the **bias-variance trade-off** in mind! We can decompose the error into bias and variance



Self-Test Questions

- 1 Why should you split the data into train, dev and test sets?
- 2 You perform 10-fold cross validation. How many models do you have to learn? Which one do you use in production?
- 3 What is the problem with accuracy?
- 4 Why do we apply the harmonic mean to compute the F1 score?
- 5 Your model gets an AUC value of 0. What does this mean?
- 6 Random search is usually preferred to optimize hyper-parameters. Why?
- 7 Your model does not perform well due to its high bias. Your boss suggests adding more training examples. How would you respond?

What's next...?

Unit I	Machine Learning Introduction
Unit II	Mathematical Foundations
Unit III	Bayesian Decision Theory
Unit IV	Regression
Unit V	Classification I
Unit VI	Evaluation
Unit VII	Classification II
Unit VIII	Clustering
Unit IX	Dimensionality Reduction

Thank you very much for the attention!

Topic: *** Applied Machine Learning Fundamentals *** Evaluation of ML Models

Term: Winter term 2023/2024

Contact:

Daniel Wehner, M.Sc.

SAP SE / DHBW Mannheim

daniel.wehner@sap.com

Do you have any questions?